

mmpdb 3.0

Andrew Dalke <dalke@dalkescientific.com>
RDKit User Group Meeting
13 October 2022

Summary

Two molecules form a matched molecular pair (MMP) if they only differ by a single set of connected atoms.

**The connected atoms are the "variable" parts.
The rest are the "constant".**

variable1>>variable2 defines a transformation.

**When chemistry is additive, transformations can
be used to predict property changes.**



**Including the local environment of the constant
can help improve linearity.**

In mmpdb 3.0:

- better scaling for large databases**
- can use MMPs to generate new structures based
on known chemical space (no AI here!)**
- environment information based on Morgan SMARTS**

Uses the "fragment-and-index" approach of Hussain and Rea



1-cut pair (mmpdb has special support for a single hydrogen)



Canonicalize the constant then find matching constants.

	Constant	Variable
Compound A	<chem>[*:1]N(C)C(=O)c1ccccc1</chem>	<chem>[*:1][H]</chem>
Compound B	<chem>[*:1]N(C)C(=O)c1ccccc1</chem>	<chem>[*:1]C</chem>

2-cut fragmentation



	Constant	Variable
Compound A	<chem>C[*:2].c1ccc([*:1])cc1</chem>	<chem>O=C(N[*:2])[*:1]</chem>
Compound B	<chem>C[*:2].c1ccc([*:1])cc1</chem>	<chem>CN(C(=O)[*:1])[*:2]</chem>

Two components in the constant.

Complications

- Can't store the attachment points as labeled wildcards ("`[*:1]`", "`[*:2]`", "`[*:3]`") because the value affects canonicalization order.
- Attachment points may be in the same symmetry class.
- Fragmentation can create symmetry causing loss of chiral information. ("`F[C@](Cl)(Br)O`" -> "`F[C@](*)(*)*`")
 - Which you may or may not want when indexing.
- The attachment point order in the constants may be different than the attachment point order in the variables.

Took a lot of work to handle these correctly!

"mmpdb smifrag"

```
% mmpdb smifrag 'CN(C)C(=O)c1ccccc1'
```


		variable				constant			
#cuts	enum.label	#heavies	symm.class	smiles	order	#heavies	symm.class	smiles	with-H
1	N	10	1	*N(C)C(=O)c1ccccc1	0	1	1	*C	C
1	N	1	1	*C	0	10	1	*N(C)C(=O)c1ccccc1	CNC(=O)c1ccccc1
2	N	9	11	*N(*)C(=O)c1ccccc1	01	2	11	*C.*C	-
→ 3	N	3	122	*C(=O)N(*)*	201	8	112	*C.*C.*c1ccccc1	-
2	N	4	12	*C(=O)N(*)C	10	7	12	*C.*c1ccccc1	-
1	N	6	1	*c1ccccc1	0	5	1	*C(=O)N(C)C	CN(C)C=O
1	N	5	1	*C(=O)N(C)C	0	6	1	*c1ccccc1	c1ccccc1

Variable

C(=O)N()*

Symmetry class "122".

Last two "*" are in the same class.



Constant

*C.*C.*c1ccccc1

Symmetry class "122".

First two "*" are in the same class.

Order "201" describes how to re-connect the constant and variable.


Wildcard attachments: [*:2]C(=O)N([*:0])[*:1].[*:0]C.[*:1]C.[*:2]c1ccccc1

As closures: C%92(=O)N%90%91.C%90.C%91.c%921ccccc1

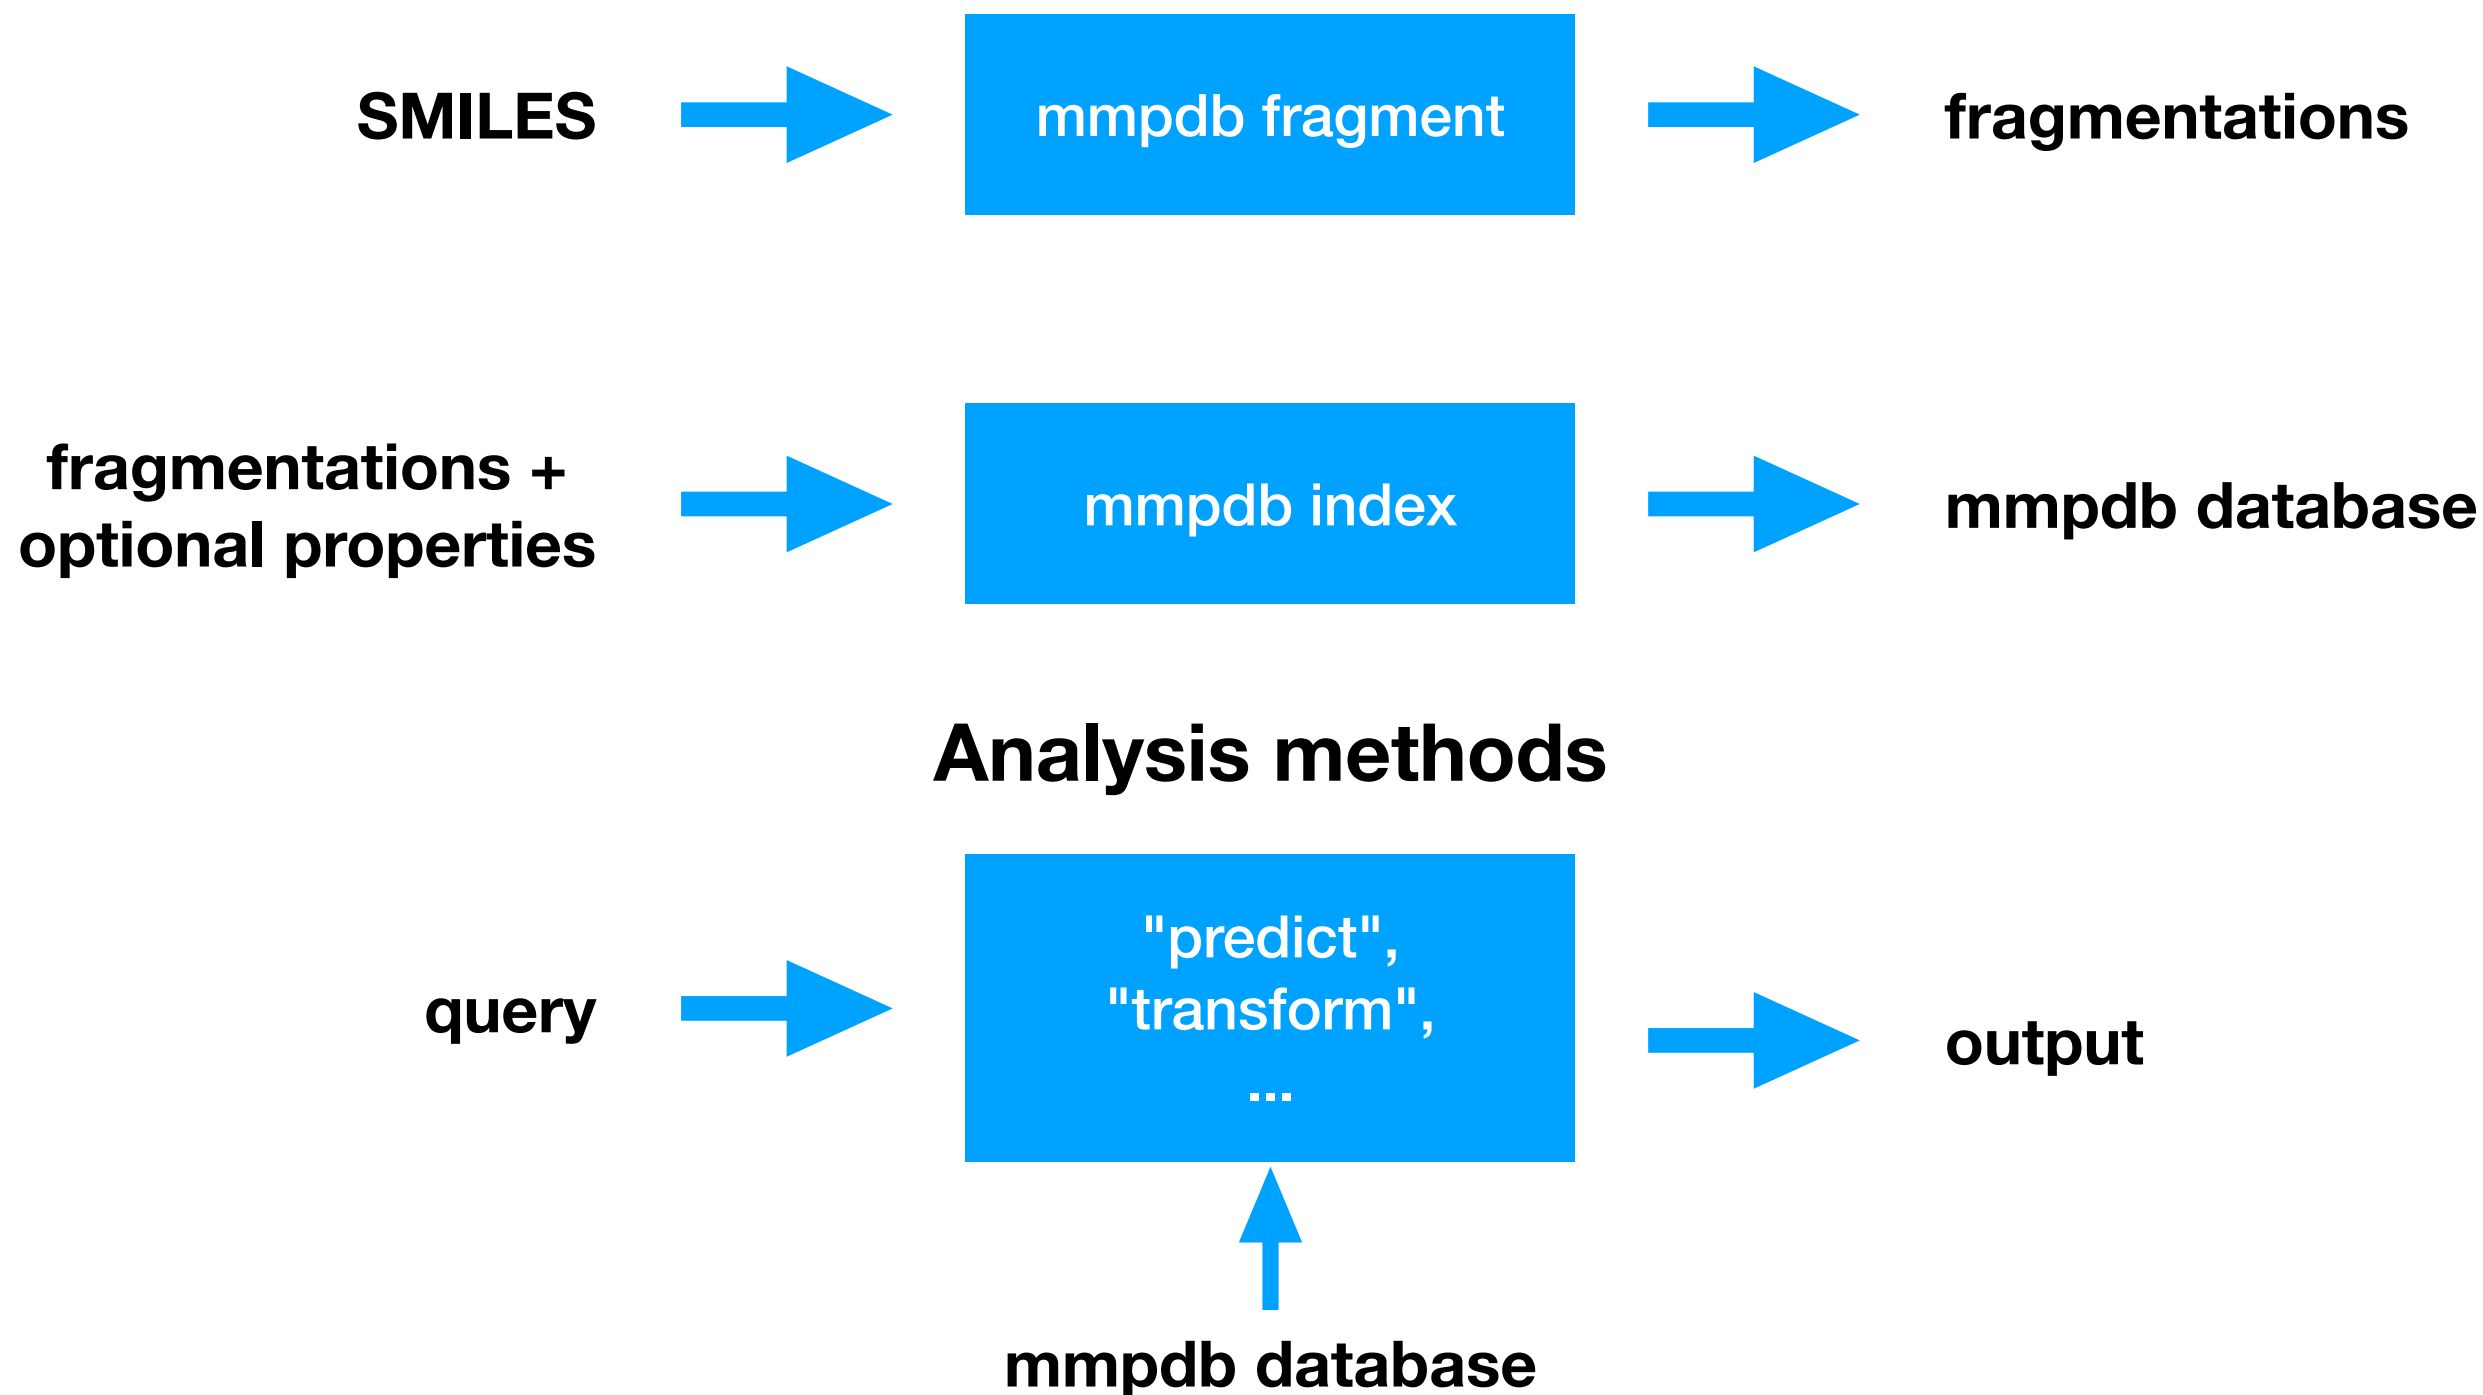
"Welded": CN(C)C(=O)c1ccccc1

Environment

"MMP rules are highly dependent on the local environment around transformations. A transformation that substitutes a hydrogen atom in a carboxylic acid with a methyl group, for example, results in different molecular property changes than the same substitution in an aliphatic chain."



Basic mmpdb Data Flow



"predict"

Predict the effect of substituting a sulfur in diphenyl ether if the known melting point is 12C.

```
% mmpdb predict csd.mmpdb --smiles c1ccccc1Sc1ccccc1 \  
--reference c1ccccc1Oc1ccccc1 --property MP --value 12.0  
predicted delta: +4.91667 predicted value: 16.9167 +/- 18.0477
```

**Can also save details about the transformations
and associated pairs to files.**

"transform"

Generate all of the products of diphenyl ether using the MMP transforms where there are at least 30 pairs. Also include the predicted effects on the 'MP' property.

```
% mmpdb transform csd.mmpdb --smiles 'c1ccccc1Oc1ccccc1' --min-pairs 30 -p MP
```

ID	SMILES	MP_from_smiles	MP_to_smiles	MP_radius	\
1	Brc1ccc(Oc2ccccc2)cc1	[*:1]c1ccccc1	[*:1]c1ccc(Br)cc1	0	
2	C0c1ccc(Oc2ccccc2)cc1	[*:1]c1ccccc1	[*:1]c1ccc(OC)cc1	0	
3	C0c1ccccc1	[*:1]c1ccccc1	[*:1]C	0	

	MP_fingerprint	MP_rule_environment_id	\
59SlQURkwt98B0D1VlKTGRkiqFDbG6JVkeTJ3ex3b0A		947	
59SlQURkwt98B0D1VlKTGRkiqFDbG6JVkeTJ3ex3b0A		4560	
59SlQURkwt98B0D1VlKTGRkiqFDbG6JVkeTJ3ex3b0A		90	

MP_count	MP_avg	MP_std	MP_kurtosis	MP_skewness	MP_min	MP_q1	\
34	14.5290	30.990	-0.267780	0.32663	-66	-7.0	
56	8.7143	38.945	7.013600	1.81870	-172	-10.0	
106	-23.4430	36.987	1.563800	0.65077	-159	-44.0	

MP_median	MP_q3	MP_max	MP_paired_t	MP_p_value
15.5	37.0	67	-2.7338	9.987200e-03
10.5	32.5	79	-1.6745	9.971500e-02
-20.0	-3.0	49	6.5256	2.447100e-09

More on this fingerprint coming up.

mmpdb 3.0

fragmentation format

Previously the fragmentations were in JSON-Lines format.

- Large text file (compresses well).
- Needed 3rd party JSON parser for fast parsing.
- Hard to re-use the fragmentations for other purposes.

Switched to SQLite

```
% sqlite3 ChEMBL_CYP3A4_hERG.fragdb
SQLite version 3.38.5 2022-05-06 15:25:27
Enter ".help" for usage hints.
sqlite> .mode line
sqlite> SELECT id FROM record WHERE normalized_smiles = 'C0c1ccc(C)cc1N';
id = 12302
sqlite> SELECT attachment_order, variable_smiles, constant_smiles
...> FROM fragmentation
...> WHERE record_id = 12302 AND num_cuts = 3;
attachment_order = 012
variable_smiles = *0c1ccc(*)cc1*
constant_smiles = *C.*C.*N

attachment_order = 021
variable_smiles = *c1ccc(*)c(*)c1
constant_smiles = *C.*N.*0C
```

With a bit of SQL, can generate a Free-Wilson table. (Still tricky to handle symmetry and chirality. And hydrogens.)

fragdb_list

Summarize the contents of a fragdb file.

```
% mmpdb fragdb_list ChEMBL_CYP3A4_hERG.*.fragdb
```

Name	#recs	#errs	#frags	#consts	#vars	max.#pairs
ChEMBL_CYP3A4_hERG.0000.fragdb	1929	97	109087	48895	66808	13267833
ChEMBL_CYP3A4_hERG.0001.fragdb	1780	246	212777	70915	146526	116930089
ChEMBL_CYP3A4_hERG.0002.fragdb	1999	27	100594	52370	57399	9958159
ChEMBL_CYP3A4_hERG.0003.fragdb	1690	336	103350	49021	67210	8544224
ChEMBL_CYP3A4_hERG.0004.fragdb	1919	107	112930	52318	68386	8580341
ChEMBL_CYP3A4_hERG.0005.fragdb	1783	243	123463	55977	72733	15043504
ChEMBL_CYP3A4_hERG.0006.fragdb	1869	157	164259	64083	106427	52277090
ChEMBL_CYP3A4_hERG.0007.fragdb	1862	164	114113	51605	61439	15083149
ChEMBL_CYP3A4_hERG.0008.fragdb	1989	37	80613	44149	41764	4614913
ChEMBL_CYP3A4_hERG.0009.fragdb	1939	87	69029	35889	36577	2919084

fragdb_constants

Show counts for all (or selected) constants.

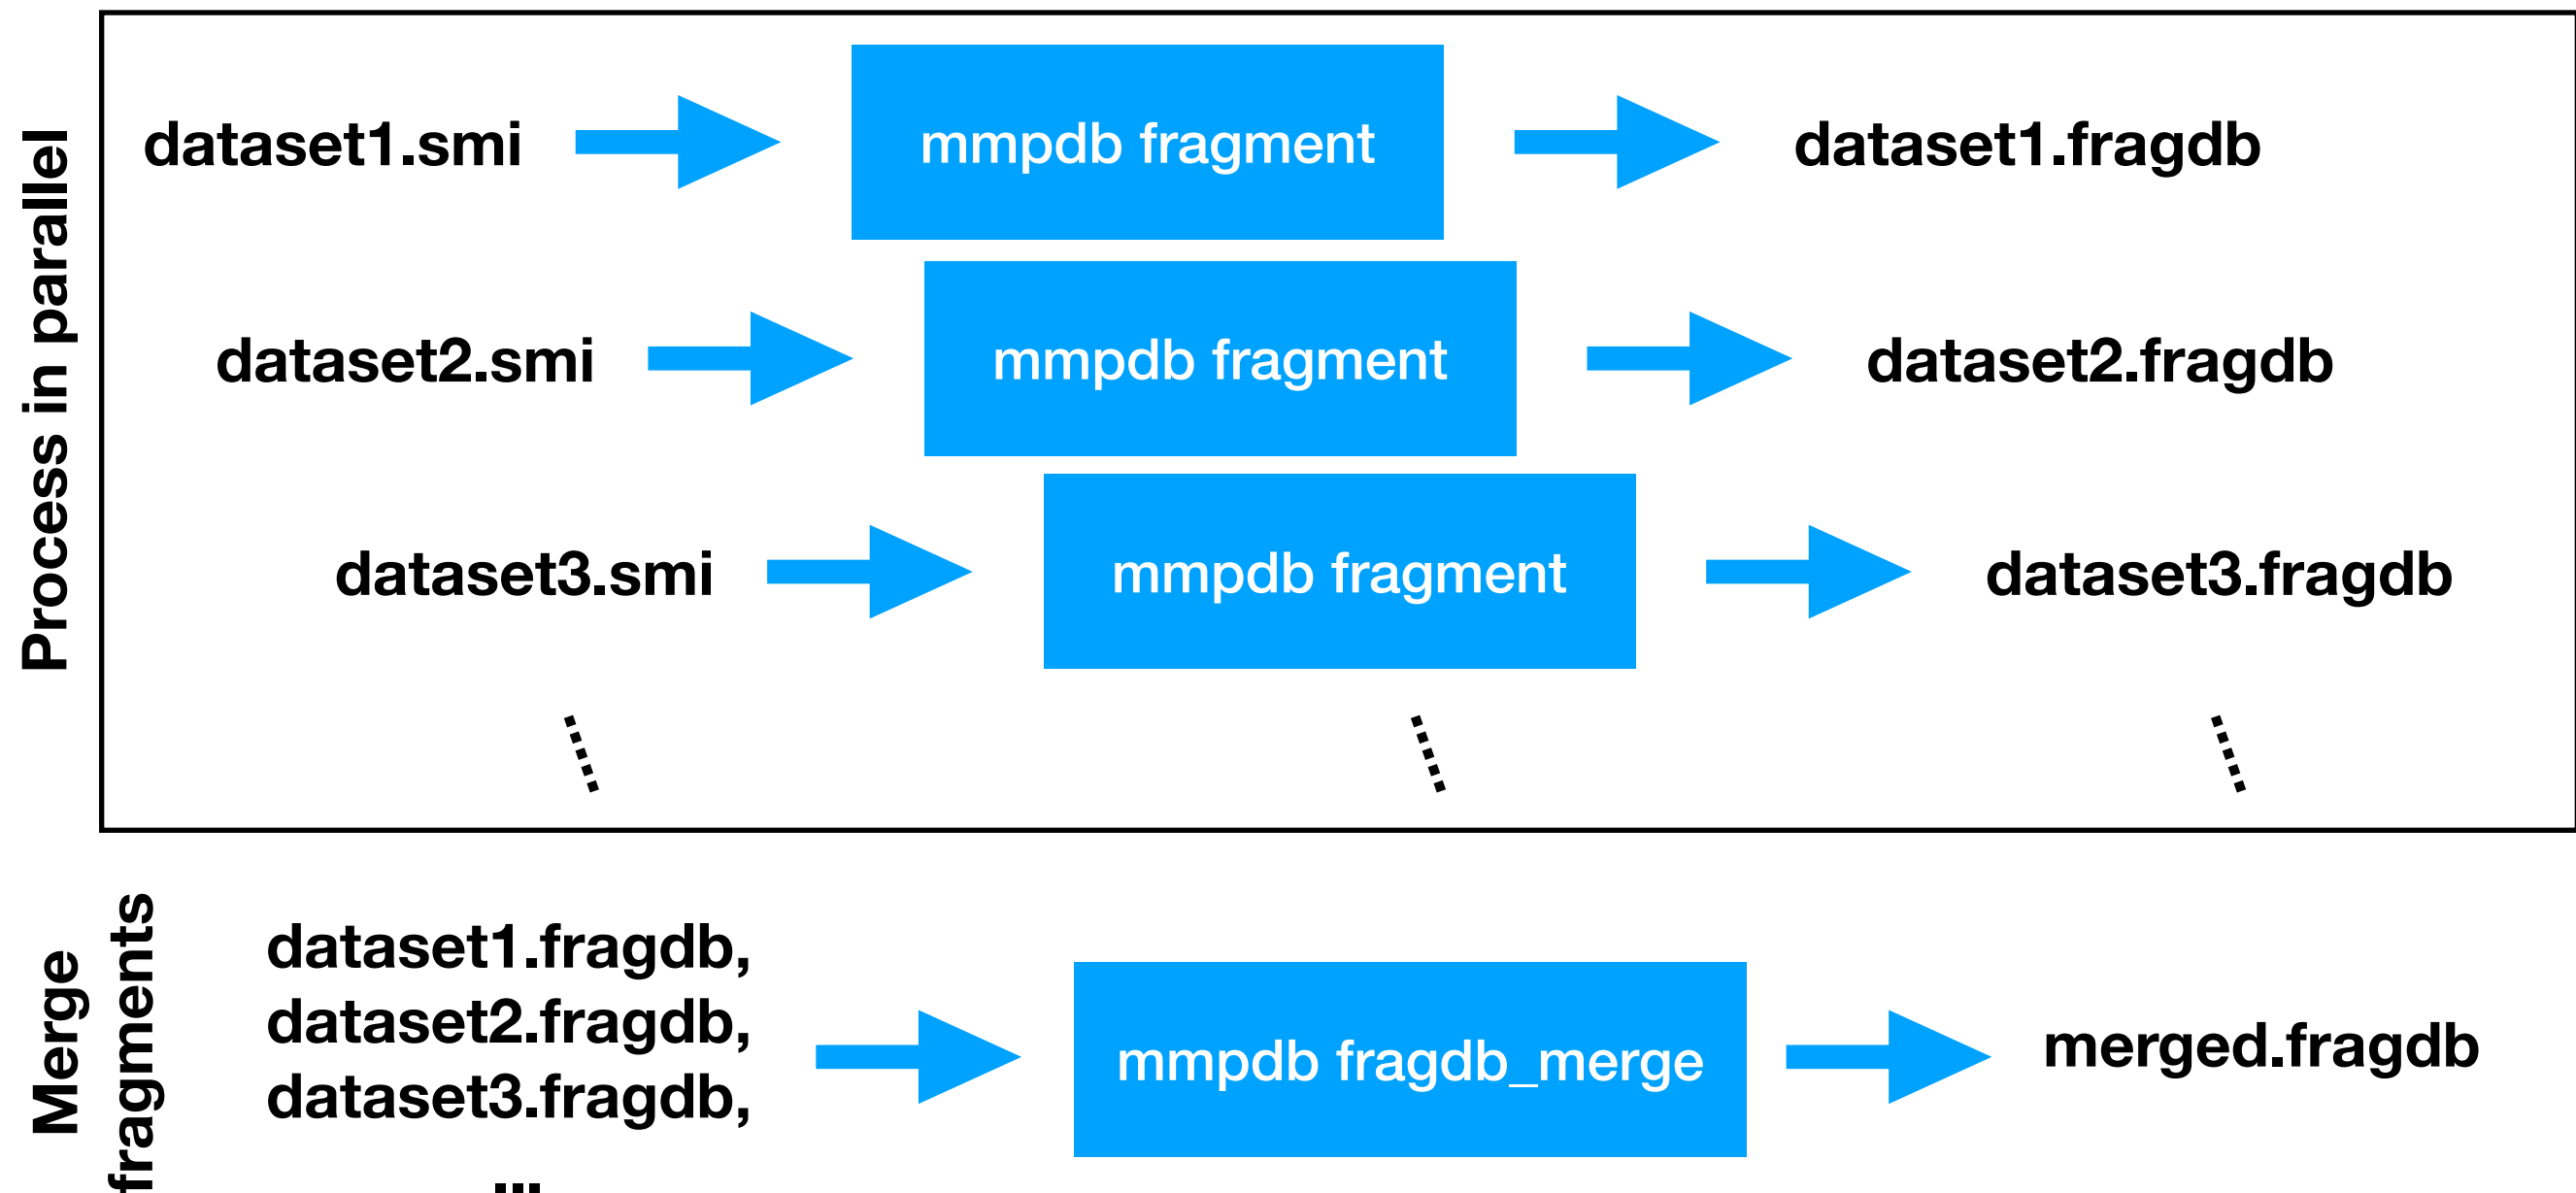
```
% mmpdb fragdb_constants test.fragdb --min-heavies-per-const-frag 5 --limit 10
constant  N
*c1cccc1      304
*N1CCC[C@H]1C  90
*c1cccc1.*c1cccc1  90
*CCN1CCC[C@H]1C  86
*c1cccc1      85
*c1ccc2cc(CCN3CCC[C@H]3C)ccc2n1  67
*c1cccn1      53
*c1ccc(Br)cc1  44
*c1ccc([N+](=O)[O-])cc1  40
*c1ccncc1     36
```

"From this analysis, we created a list of constant components that occurred 1000 times or more in either the RocheDB or SureChEMBL databases and discarded them from further consideration. This step is important to discard spurious MMP's that otherwise exponentially increase not only the transformation database size but also the computation power and memory requirement."

Merge fragmentations

Can be used to fragment in parallel.

Only use for single-threaded indexing, or if you want a merged fragment db.



Parallel indexing

Approach developed by Mahendra Awale while at Roche.*

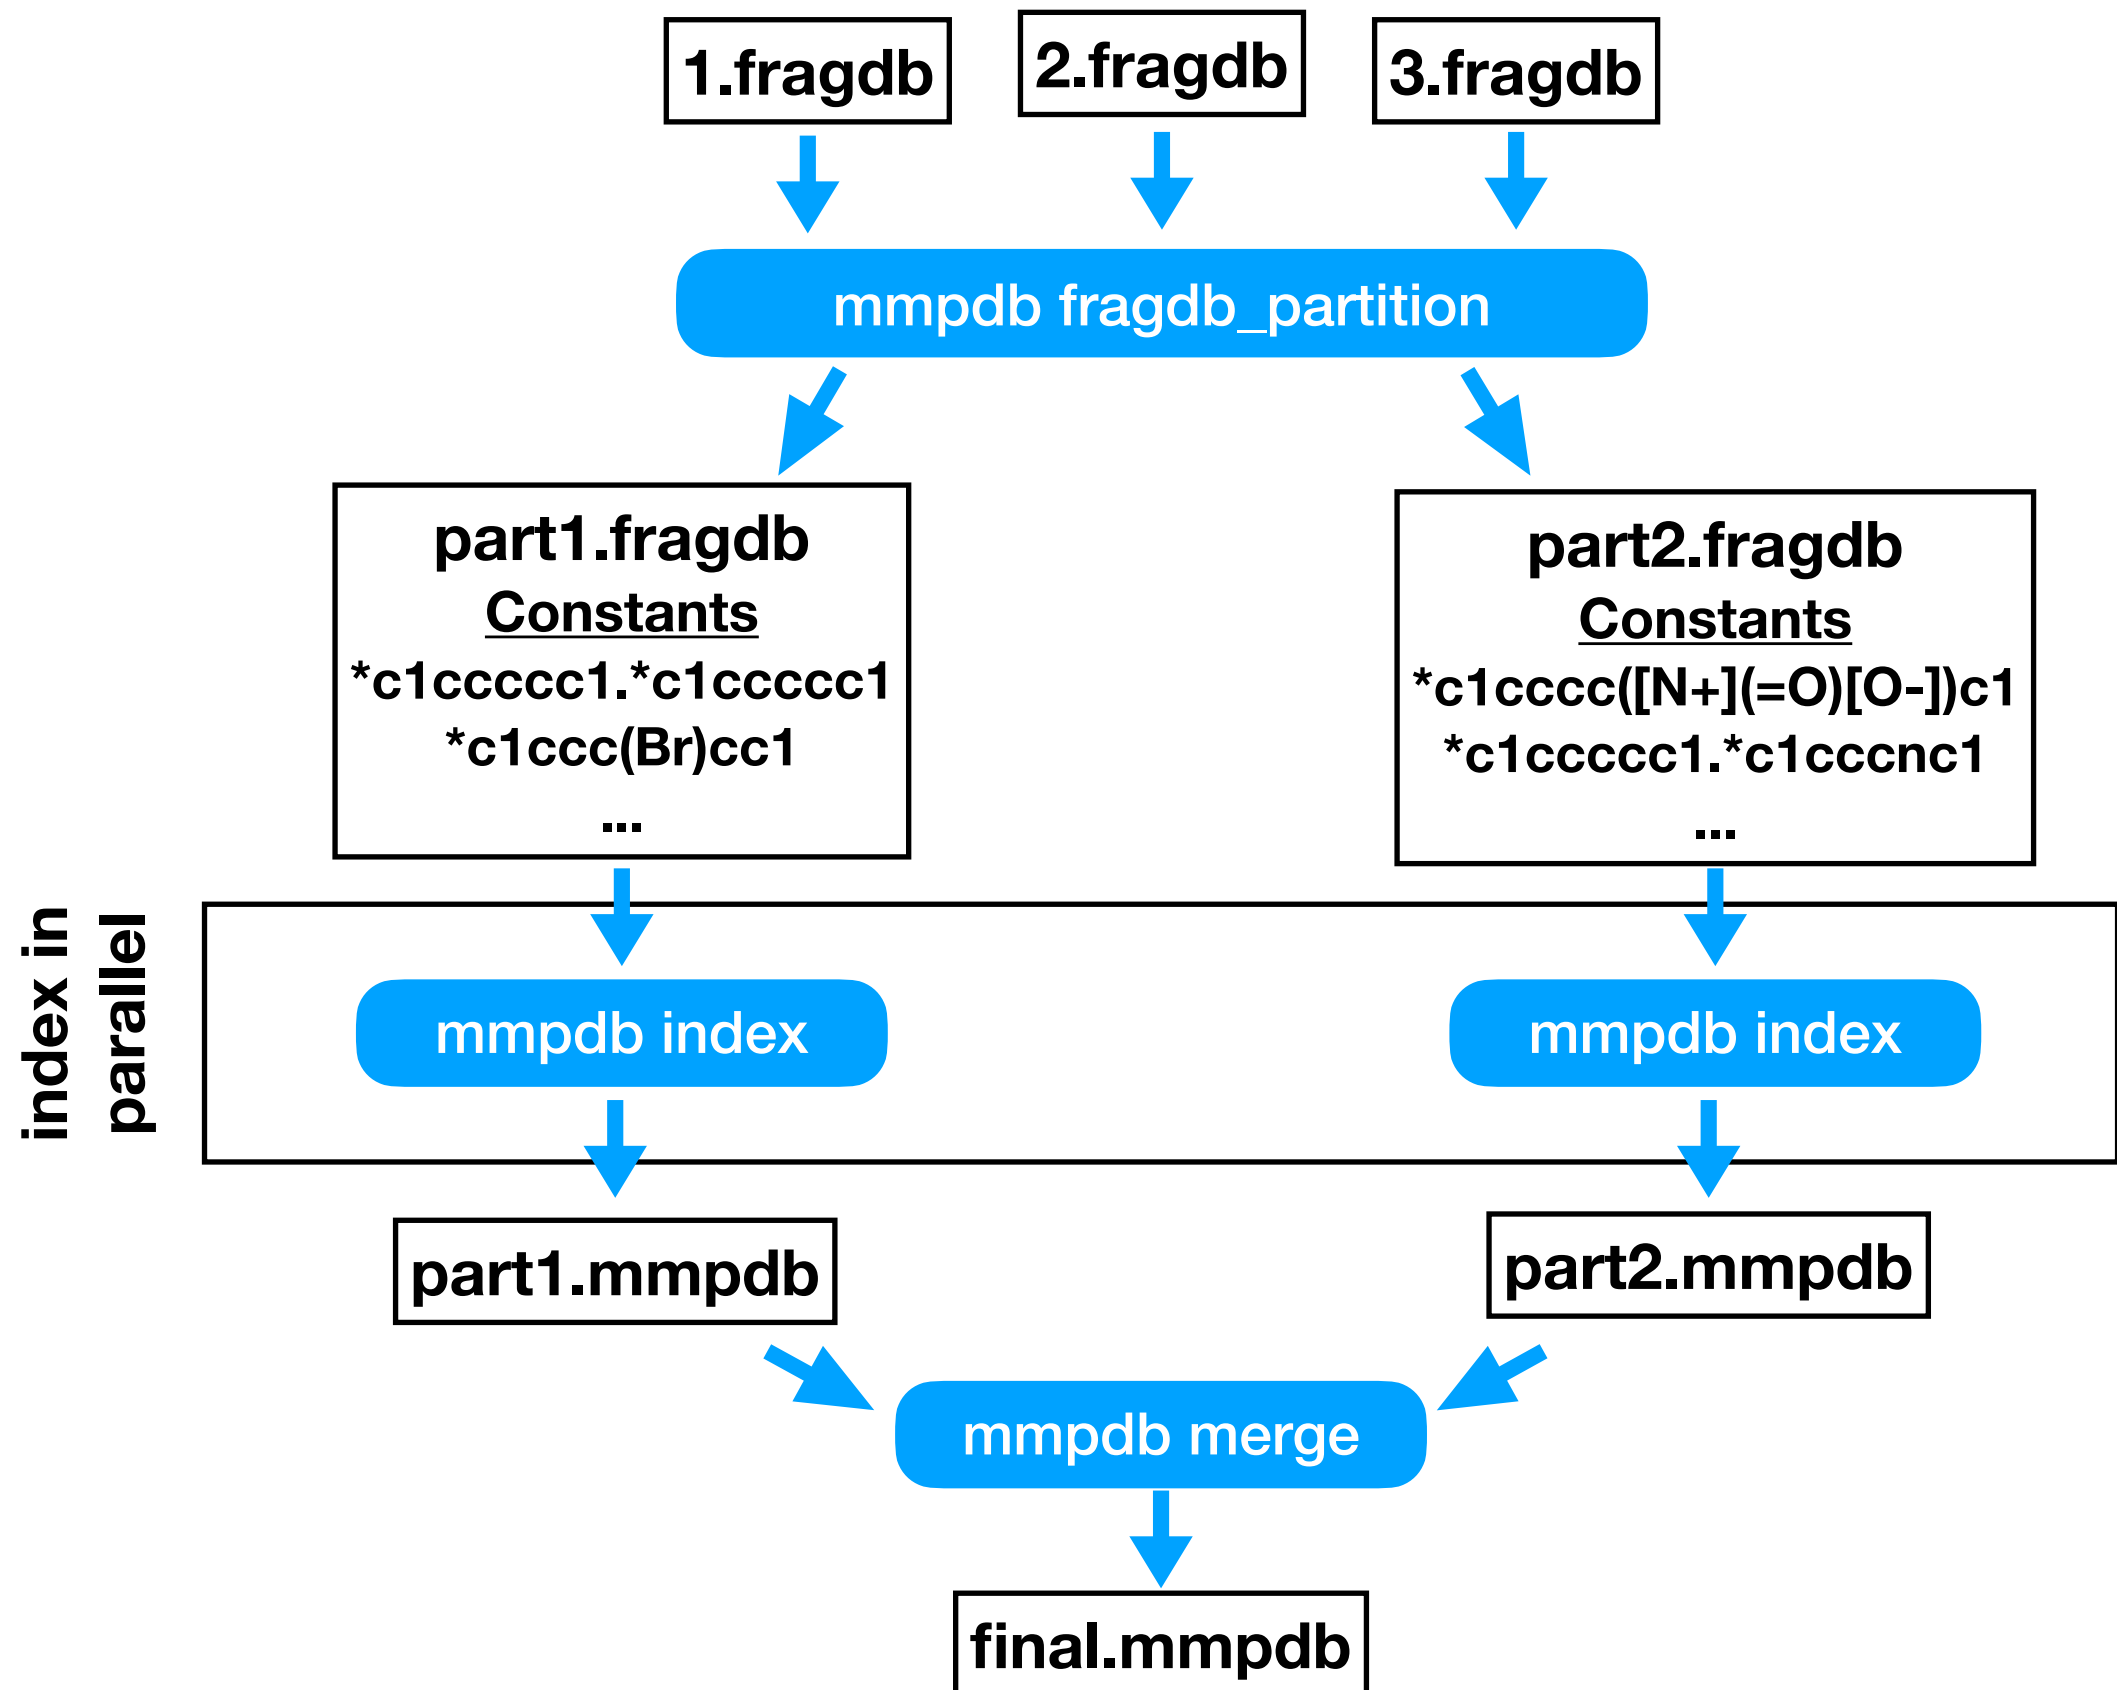
**Fragment database merging is single-threaded.
Indexing is also single-threaded ...
... *for a given constant!***

**Partition the fragmentations by constant to
generate new fragment databases.**

- Indexing cost is $\sim O(N^2)$ in the constant's number of fragmentations.**
- Round-robin to N output files, or use max cost per file.**

*** *The Playbooks of Medicinal Chemistry Design Moves*
Mahendra Awale, Jérôme Hert, Laura Guasch, Sereina Riniker, and Christian Kramer
<https://pubs.acs.org/doi/10.1021/acs.jcim.0c01143>**

Parallel data flow



Attach SQLite databases

"merge" is still single-threaded.

Needs to normalize a lot of tables and fields.

Current system uses SQLite's ability to "attach" a database.

All connected through an SQLite :memory: database.

```
working_db = sqlite3.connect(":memory:")  
schema.create_schema_for_sqlite(working_db)  
c = working_db.cursor()
```

Store mappings in an
in-memory database.

```
c.execute("ATTACH DATABASE ? AS new", (output_filename,))
```

Attach the output
database.

```
...  
process_compound_tables(c, filenames, reporter)  
process_rule_smiles_tables(c, filenames, reporter)  
process_rule_tables(c, filenames, reporter)  
process_environment_fingerprint_tables(c, filenames, reporter)  
process_rule_environment_tables(c, filenames, reporter)  
process_pair_tables(c, filenames, reporter)  
...
```

Each "process_*()" uses:

```
c.execute("ATTACH DATABASE ? AS old", (filename,))  
...  
c.execute("DETACH DATABASE old")
```

Attach then process
the input database.

Example merge SQL

```
def process_pair_table(c, db_id):
    c.execute("""
INSERT INTO new.pair (rule_environment_id, compound1_id, compound2_id, constant_id)
    SELECT rule_environment_map.new_rule_environment_id,
           compound1_map.new_compound_id,
           compound2_map.new_compound_id,
           new_constant_smiles.id

FROM old.pair AS old_pair,
     rule_environment_map_{db_id} AS rule_environment_map,
     compound_map_{db_id} AS compound1_map,
     compound_map_{db_id} AS compound2_map,
     old.constant_smiles as old_constant_smiles,
     constant_smiles as new_constant_smiles
WHERE old_pair.rule_environment_id = rule_environment_map.old_rule_environment_id
      AND old_pair.compound1_id = compound1_map.old_compound_id
      AND old_pair.compound2_id = compound2_map.old_compound_id
      AND old_pair.constant_id = old_constant_smiles.id
      AND old_constant_smiles.smiles = new_constant_smiles.smiles
;
""").format(db_id=db_id))
```

**Advice: Don't treat SQLite as simple table storage.
Do more data processing in SQL instead of Python.**

Timing info

ChEMBL 25 @ ~1.6M compounds

smi_split: ~30s (into 250 files)

fragment: ~2300s+/-800s (per task, 250 files in parallel)

frag_constants + frag_partition dry run: 125s

frag_partition: ~88s+/-7s (per task, 31 files in parallel)

index: ~209s+/-275s (per task, 31 files in parallel;

last 2 tasks took 800s and 1500s)

merge: ~2000s

resulting file: 25GB

**~1 CPU week
(serial)**

**~ 2 hours
(parallel)**

SureChEMBL (Oct. 2019) @ ~13.5M compounds

smi_split: 140s

fragment: ~7000s+/-1700s (per task, 250 files in parallel)

frag_constants + frag_partition dry run: 1124s

frag_partition: ~720s+/-220s (per task, 56 files in parallel)

index: ~1377s+/-1952s (per task, 56 files in parallel;

last 2 tasks took 9200s and 11500s)

merge: ~24 hours (requires a lot of memory! 512GB node)

resulting file: ~200GB

**~3 CPU week
(serial)**

**~ 30 hours
(parallel)**

New "generate" method

The MMP database can be seen as a medchem "playbook".*
Convert implicit database information into explicit design knowledge.
Apply those rules to a structure to generate new structures.

Similar to "transform" but with more control
over which part to substitute.
Only works for 1-cut fragmentations.

**Specify only
a SMILES**

Generate all
transforms using all
fragmentations.

**Specify constant
and query**

Generate transforms
only for that query.

**Specify SMILES
and constant or query**

mmpdb will figure out the
other component, then
generate transforms for
that query.

** The Playbooks of Medicinal Chemistry Design Moves*

Mahendra Awale, Jérôme Hert, Laura Guasch, Sereina Riniker, and Christian Kramer


<https://pubs.acs.org/doi/10.1021/acs.jcim.0c01143>

Generate from aspirin


```
% mmpdb generate --smiles 'O=C(C)Oc1ccccc1C(=O)O' --query '*OC(C)=O'\n--radius 1 merged.mmpdb --columns #pairs,final
```

```
#pairs final
```


```
4 O=C(O)c1ccccc1O
1 COc1ccccc1C(=O)O
1 CCN(CC)c1ccccc1C(=O)O
1 NNCc1ccccc1C(=O)O
1 O=C(O)c1ccccc1Nc1ccccc1
1 O=C(Oc1ccccc1C(=O)O)c1cccs1
```




aspirin




#1




#2




#3



#4



#5




#6

Additionally, can use `--subqueries` to generate subfragments of the query fragments and use those as additional query fragments.

Environment Fingerprint

Require a transform rule like [*:2]NC([*:1])=O>>[*:1]C(=O)N([*:2])C to also match some of the circular environment around the attachment points.

Constant environment
with radius 2



Old method

- 1) Compute Morgan fingerprints centered at each attachment point.
- 2) `hashlib.sha256(fp.ToBinary()).digest()`
- 3) Concatenate the sha256s.
- 4) SHA256 the concatenation.
- 5) Base64-encode the result.

59SIQURkWt98BOD1VIKTGRkiqFDbG6JVkeTJ3ex3bOA

Only useful for identify match. Fixed-width size. Uninterpretable.

Morgan Atom SMARTS

The Morgan atom connectivity invariants are:

```
components.push_back(atom->getAtomicNum());  
components.push_back(atom->getTotalDegree());  
components.push_back(atom->getTotalNumHs());  
components.push_back(atom->getFormalCharge());
```

.. more about mass, but mmpdb removes isotopes on input ...

```
if (includeRingMembership &&  
    atom->getOwningMol().getRingInfo()  
        ->numAtomRings(atom->getIdx())) {  
    components.push_back(1);  
}
```

These can be represented directly in SMARTS:

[#7,X3,H2,+1,R]

Morgan Bond SMARTS

The Morgan bond connectivity invariant is:

```
bondInvariant = static_cast<int32_t>(bond->getBondType());
```

The mapping to SMARTS is:

```
_bond_smarts_symbols = {  
    Chem.BondType.SINGLE: "-",  
    Chem.BondType.DOUBLE: "=",  
    Chem.BondType.TRIPLE: "#",  
    Chem.BondType.AROMATIC: "~",  
}
```

```
... _bond_smarts_symbols[bond.GetBondType()] ...
```

Fragment SMARTS

**MolFragmentToSmiles lets you specify the atom and bond symbols.
Use them to generate a SMARTS.**

```
smarts = Chem.MolFragmentToSmiles(  
    mol,  
    atomsToUse = list(atom_ids),  
    bondsToUse = list(bond_ids),  
    atomSymbols = atom_symbols,  
    bondSymbols = bond_symbols,  
    isomericSmiles = False,  
)
```

Can't specify the atom&bond invariants.

The output will likely not be canonical ...

... unless the atom and bond symbols match the internal invariants!

Disable isomeric SMILES because Morgan fingerprints don't use them.

Canonical SMARTS for a Rooted Morgan Fingerprint


- 1) Get the atom and bond symbols.
- 2) Set aromatic atoms to aliphatic.
(To match the Morgan invariants)
- 3) `Chem.MolFragmentToSmiles()`
- 4) Reorder SMILES components;
First `[:1]` then `[:2]` and `[:3]`
- 5) Restore aromatic atoms.

Only works because the SMARTS is rooted at the attachment point.

Morgan SMARTS

Can be quite large!
(VARCHAR(1024) instead of 32)

**[#0;X1;H0;+0;!R:1]-[C;X4;H2;+0;!R]-[O;X2;H0;+0;!R]-[#6;X3;H0;+0;R]
(:[#6;X3;H1;+0;R]:[#6;X3;H0;+0;R]):[#6;X3;H1;+0;R]:[#6;X3;H1;+0;R]**



SMARTSViewer doesn't accept #0 so I used #100

Picture created by the SMARTviewer [<https://smarts.plus/>].
Copyright: ZBH - Center for Bioinformatics Hamburg.


pseudo-SMILES

Most people can't easily read SMARTS.
Few tools can visualize a complex SMARTS.

Convert the SMARTS into a SMILES that RDKit will parse with sanitize=False.

[*:1]-[CH2]-[O]-[c](:[cH]:[c](~*)(~*)):[cH]:[cH](~*)

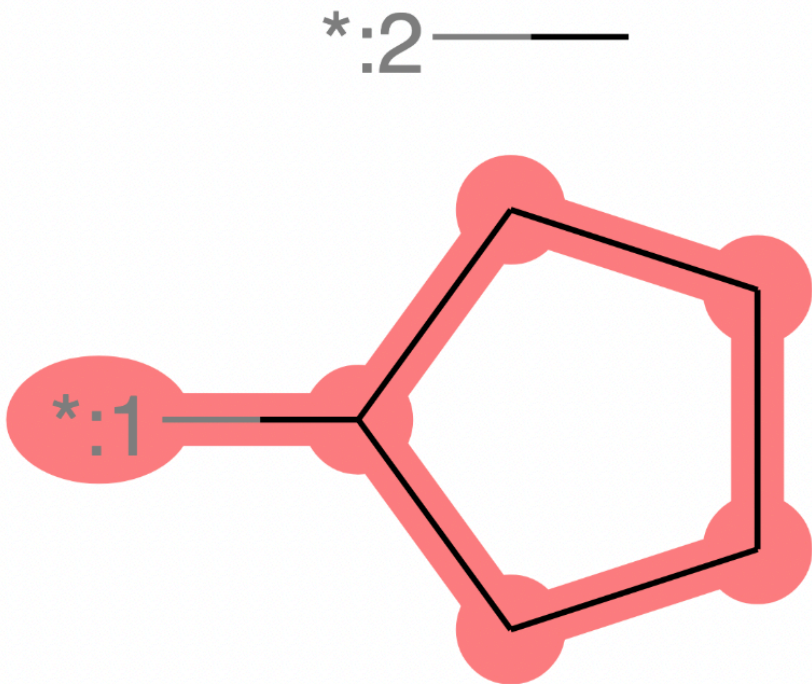
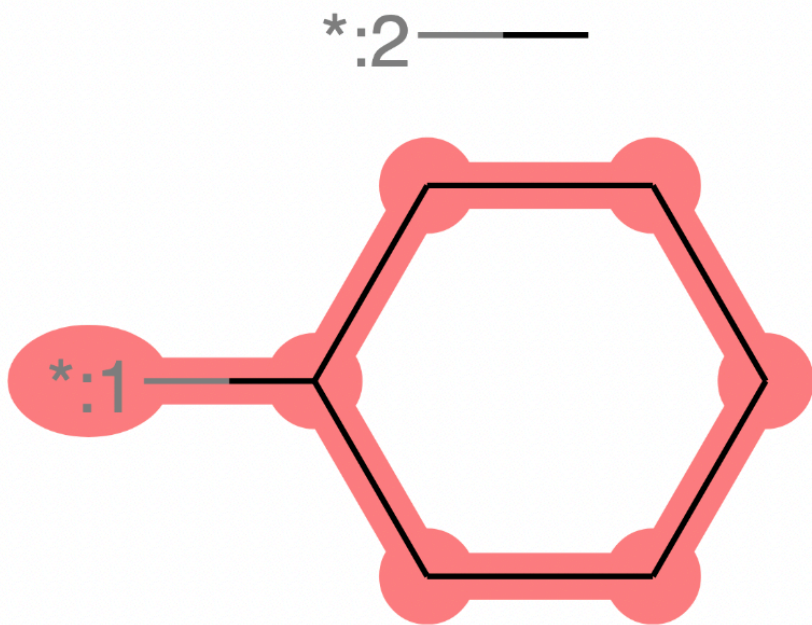
[O;X2;H0;+0;!R] **↑** **Needed for the X3.** **↑** **[:#6;X3;H0;+0;R]**



Watch out! Two distinct "*" atoms in the pseudoSMILES *might* refer to the same atom in the original molecule!

Cross-checking

***Almost* perfect one-to-one mapping from
SHA256-hashed fingerprint to Morgan SMARTS.
Can have multiple SMARTS for the same SHA256 fp.**

fingerprint Tcu44H+ojDtSVcNYf9pQmSOIAOtxfmCKHzfQZGWEzAs	
env #1 (r = 4)	env #2 (r = 4)
context: C1CCC([*:1])C1.C[:2]	context: C1CCC([*:1])CC1.C[:2]
smarts: [#6&X4&H2&+0&R]1-[#6&X4&H2&+0&R]-[#6&X4&H2&+0&R]-[#6&X4&H1&+0&R](-[#6&X4&H2&+0&R]-1)-[#0&X1&H0&+0&!R:1]	smarts: [#6&X4&H2&+0&R]1-[#6&X4&H2&+0&R]-[#6&X4&H2&+0&R]-[#6&X4&H1&+0&R](-[#6&X4&H2&+0&R]-1)-[#0&X1&H0&+0&!R:1]
	

Other changes

Added support for Postgres for direct index database creation.

Does not handle the new parallelized indexing.

(Use one of the SQLite→Postgres converters instead?)

Also added a "csvd" output.

Save tables to a directory containing CSV files.

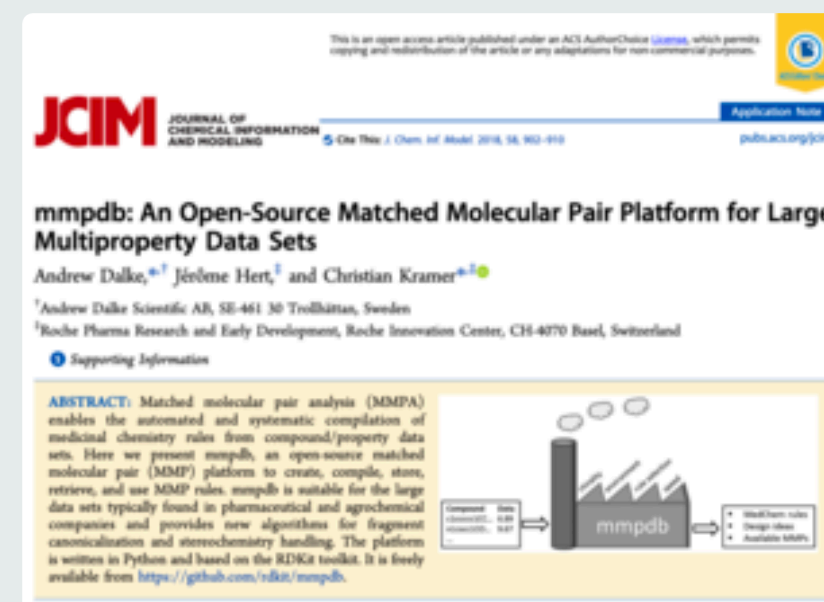
'mmpdb proprulecat' command to export the property rules in the database (transformation + statistics) in CSV form

Moved from argparse to click.

mmpdb crowdfunding project

How can we **raise money to fund open source software development in cheminformatics**? It's a hard question.

Simple donations don't work – companies might not even have a mechanism to make donations. **Consultant-based funding doesn't work** that well either, because the cost of developing a general-purpose tool is two or three times more expensive than developing a tool which only meets the specialized needs of one client, and few clients are willing to subsidize the rest of the field. Proprietary software development solves the problem by getting many people to pay for the same product. **Can we learn from the success of proprietary software** to get the funds which would certainly be useful in improving open source software?



J. Chem. Inf. Model. 2018, 58, 902–910.

ACS Editors' Choice

Consortium model

- **My company is the supplier/organizer.**
 - **Easy to invoice.**
- **Anyone could pay to join. Suggested prices:**
 - **Academics - EUR 1 000 (no warranty, limited support)**
 - **Industry - EUR 5 000 (includes 9 months of support)**
- **Members get the source code under an open source license.**
 - **No promise the resulting work would be made public.**
- **Started the effort *after* I had some new features in place.**
 - **A company paid me (as a consultant) to add them.**
 - **I asked for permission to use that code as a seed.**
 - **Could promise those features to all consortium members.**
 - **Made them the first consortium member.**

Funding Levels and Goals

**Companies are more willing to fund features.
Need to improve mmpdb's "infrastructure."**

**Raised
EUR 22 500**



EUR 16 000 - Environment fragment SMILES

EUR 23 000 - Public release after 9 months

EUR 29 000 - Documentation

EUR 34 000 - mmpdb/GitHub user support

EUR 40 000 - Test suite development

EUR 50 000 - Immediate public release

**Not enough
funding for
infrastructure
improvements.**



These prices are low compared to industry standards!

Plus overhead for marketing, web designer, and accounting.

Roche Funding

Roche funded nearly all of mmpdb development.

**They had an in-house branch to handle large data sets,
and they implemented the "generate" functionality.**

**Paid me to turn "research" code into "production" code.
And they were consortium members.**


**I merged the three code bases together, and
decided to release the result to the public.**

Current Status

mmpdb 3.0b1 is ready

<https://github.com/adalke/mmpdb/tree/v3-dev>


**"v3-dev"
branch** →

 v3-dev ▾

Go to file

Code ▾

This branch is [289 commits ahead](#) of rdkit:master.

 Contribute ▾

**Working on a speedup for "generate" for large DBs.
Then will push to main branch under RDKit's account.**

Future

No long-term development or support plans.

Christian Kramer @ Roche provides support in his spare time.

**Still missing infrastructure (documentation and testing)
I think these are important for future development.**

I don't know how to get stable long-term funding.

Thanks!

Roche

Mahendra Awale

Jérôme Hert

Christian Kramer

and to all the members of the mmpdb consortium!